

Ensemble

Big Data Lectures – Chapter 6

Lexin Li

Division of Biostatistics
University of California, Berkeley



Outline

- ▶ list of topics:
 - ▶ neural networks and deep learning
 - ▶ boosting
 - ▶ random forest
 - ▶ more ensemble methods



Neural Networks and Deep Learning




Neural networks

- ▶ neural networks:
 - ▶ the term has evolved to encompass a large class of models and learning methods, especially popular in the field of **artificial intelligence, computer vision, and automatic speech recognition**
 - ▶ we describe a widely used neural net, sometimes called the **single hidden layer back-propagation network**, or **single layer perceptron**
- ▶ examples:
 - ▶ a typical neural network used 1 to 10 million connections
 - ▶ Google Brain neural network used more than 1 billion connections developed a distributed computing infrastructure and spread the computation across 16,000 of our CPU cores



Neural networks

- key idea: **sum of nonlinear functions of linear combinations of the inputs**, typically represented by a **network diagram**



Neural networks

- ▶ neural networks model:

$$Z_m = \sigma(\alpha_{0m} + \alpha_m^T \mathbf{X}), \quad m = 1, \dots, M$$

$$T_k = \beta_{0k} + \beta_k^T \mathbf{Z}, \quad k = 1, \dots, K$$

$$f_k(\mathbf{X}) = g_k(T), \quad k = 1, \dots, K$$


- ▶ $\mathbf{Y} = (Y_1, \dots, Y_K)$ are the K -dimensional output; e.g., for univariate response, $K = 1$; for K -class classification, k th unit models the probability of class k
- ▶ $\mathbf{X} = (X_1, \dots, X_p)$ are the p -dimensional input features
- ▶ $\mathbf{Z} = (Z_1, \dots, Z_M)$ are the **derived features** created from linear combinations of the inputs \mathbf{X}
- ▶ $\mathbf{T} = (T_1, \dots, T_K)$ are the output features that are directly associated with the outputs \mathbf{Y} through output functions $g_k(\cdot)$
- ▶ $g_k(T) = T$ for regression; $g_k(T) = e^{T_k} / \sum_{l=1}^K e^{T_l}$ for K -class classification



Neural networks

► activation function:

- usually $\sigma(v) = \text{sigmoid} = 1/(1 + e^{-v})$
- $\sigma(v)$ = a step function: human brain models where each unit represents a neuron, and the connections represent synapses; the neurons fired when the total signal passed to that unit exceeded a certain threshold



Neural networks

- loss function: residual sum of squares / deviance

$$L = \sum_{i=1}^n \sum_{k=1}^K (y_{ik} - f_k(x_i))^2, \quad L = - \sum_{i=1}^n \sum_{k=1}^K y_{ik} \log f_k(x_i)$$

- model fitting: **back propagation** – gradient descent
 - gradient descent update:

$$\beta_{km}^{(r+1)} = \beta_{km}^{(r)} - \gamma_r \sum_{i=1}^n \frac{\partial R_i}{\partial \beta_{km}} \quad \alpha_{ml}^{(r+1)} = \alpha_{ml}^{(r)} - \gamma_r \sum_{i=1}^n \frac{\partial R_i}{\partial \alpha_{ml}}$$

where $R_i = \sum_{k=1}^K (y_{ik} - f_k(x_i))^2$

- the derivatives: letting $z_{mi} = \sigma(\alpha_{0m} + \alpha_m^T x_i)$

$$\frac{\partial R_i}{\partial \beta_{km}} = -2(y_{ik} - f_k(x_i))g'_k(\beta_k^T z_i)z_{mi} \equiv \delta_{ki}z_{mi}$$

$$\frac{\partial R_i}{\partial \alpha_{ml}} = - \sum_{k=1}^K 2(y_{ik} - f_k(x_i))g'_k(\beta_k^T z_i)\beta_{km}\sigma'(\alpha_m^T x_i)x_{il} \equiv s_{mi}x_{il}$$



Neural networks

- ▶ model fitting: **back propagation** (continued)
 - ▶ δ_{ki} and s_{mi} are the “errors” from the current model at the output and hidden layer units
 - ▶ back propagation equations

$$s_{mi} = \sigma'(\alpha_m^\top x_i) \sum_{k=1}^K \beta_{km} \delta_{ki}$$

- ▶ two-pass updates:
 - ▶ forward pass: fix the current weights and compute the predicted values $\hat{f}_k(x_i)$
 - ▶ backward pass: compute the errors δ_{ki} , then back-propagate to compute the errors s_{mi}
 - ▶ update $\hat{\beta}_{km}$ and $\hat{\alpha}_{ml}$
- ▶ γ_r is the learning rate
- ▶ advantages: **simple and local** nature; each hidden unit passes and receives information only to and from units that share a connection; can be implemented efficiently on a parallel architecture computer



Neural networks

- ▶ practical issues:
 - ▶ starting values: usually starting values for weights are chosen to be random values near zero; hence the model starts out nearly linear, and becomes nonlinear as the weights increase
 - ▶ overfitting: early stopping; weight decay
 - ▶ scaling of the inputs: mean zero and standard deviation one
 - ▶ how many hidden units; how many hidden layers: guided by domain knowledge and experimentation
 - ▶ multiple minima: try with different starting values



Neural networks

- ▶ practical issues:
 - ▶ starting values: usually starting values for weights are chosen to be random values near zero; hence the model starts out nearly linear, and becomes nonlinear as the weights increase
 - ▶ overfitting: early stopping; weight decay
 - ▶ scaling of the inputs: mean zero and standard deviation one
 - ▶ how many hidden units; how many hidden layers: guided by domain knowledge and experimentation
 - ▶ multiple minima: try with different starting values
- ▶ neural networks model is a **projection pursuit** type additive model:


$$f(\mathbf{X}) = \beta_0 + \sum_{m=1}^M \beta_m \sigma(\alpha_{m0} + \alpha_m^T \mathbf{X})$$


- ▶ **view**: neural networks are not a fully automatic tool, as they are sometimes advertised; as with all statistical models, subject matter knowledge should and often be used to improve their performance




Example – handwritten digits

- ▶ ZIP code data:
 - ▶ input: 256 pixel values from 16×16 grayscale images; output: 0, 1, ..., 9 10-class classification
 - ▶ a modest experimental subset: 320 training digits and 160 testing digits





Local Connectivity



Shared Weights

Example – handwritten digits

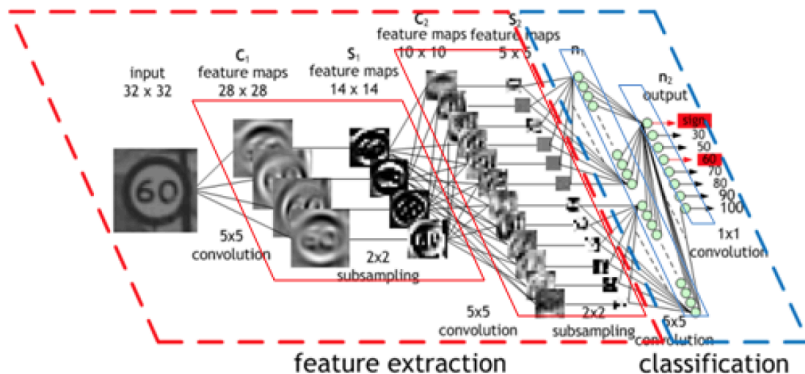
- ▶ neural networks used:
 - ▶ net-1: no hidden layer, equivalent to multinomial logistic regression
 - ▶ net-2: one hidden layer, 12 hidden units fully connected
 - ▶ net-3: two hidden layers locally connected
 - ▶ net-4: two hidden layers, locally connected with weight sharing
 - ▶ net-5: two hidden layers, locally connected, two levels of weight sharing (was the result of many person years of experimentation)
- ▶ results:

network	links	weights	accuracy
net 1	2570	2570	80.0%
net 2	3124	3214	87.0%
net 3	1226	1226	88.5%
net 4	2266	1131	94.0%
net 5	5194	1060	98.4%



Deep learning

- ▶ deep learning is a class of machine learning algorithms that:
 - ▶ use a cascade of multiple layers of nonlinear processing units, and each successive layer uses the output from the previous layer as input




Deep learning

► convolution:


- a mathematical operation that describes a rule of how to mix two functions or pieces of information
- the feature map (the input data) + the convolution kernel \Rightarrow a transformed feature map
- convolution is often interpreted as a **filter**, where the kernel filters the feature map for information of a certain kind; e.g., one kernel might filter for edges and discard other information

Input image

Convolution
Kernel

$$\begin{bmatrix} -1 & -1 & -1 \\ -1 & 8 & -1 \\ -1 & -1 & -1 \end{bmatrix}$$

Feature map



1	1 _{x1}	1 _{x0}	0 _{x1}	0
0	1 _{x0}	1 _{x1}	1 _{x0}	0
0	0 _{x1}	1 _{x0}	1 _{x1}	1
0	0	1	1	0
0	1	1	0	0


Image

4	3	

Convolved
Feature

Deep learning

- ▶ **hierarchical feature representation:**
 - ▶ form a hierarchical representation of the data, where higher level features are derived from lower level features
 - ▶ image recognition: pixel \rightarrow edge \rightarrow motif \rightarrow part \rightarrow object
 - ▶ speech: sample \rightarrow spectral band \rightarrow sound \rightarrow phoneme \rightarrow word
- ▶ feature engineering vs feature learning
 - ▶ **feature engineering:** extract useful features from the data for a specific learning task
 - ▶ **feature learning:** feature engineering **automatically done by algorithms**
— the learnt features are usually difficult to interpret



Deep learning

- ▶ challenges of a deep network:
 - ▶ **computation** — hardware advancement, e.g., graphics processing unit (GPU)
 - ▶ **vanishing gradients**: the gradients became too small to provide a learning signal for very deep layers, thus making the deep architectures perform poorly when compared to shallow learning algorithms — pre-trained one level at a time through unsupervised learning, then further fine-tuned through supervised back-propagation, alternative activation functions, alternative network structures ...
- ▶ deep architectures:
 - ▶ feed-forward: multilayer neural networks, convolutional neural networks
 - ▶ feed-back: stacked sparse coding, deconvolutional nets
 - ▶ bi-directional: deep Boltzmann machines, stacked auto-encoders



Deep learning

- ▶ learning protocols:
 - ▶ purely supervised
 - ▶ initialize parameters randomly
 - ▶ train in supervised mode using back-propagation to compute gradients
 - ▶ used in most practical systems for speech and image recognition
 - ▶ unsupervised & layerwise + supervised classifier on top
 - ▶ train each layer unsupervised, one after the other
 - ▶ train a supervised classifier on top, keeping the other layers fixed
 - ▶ good when very few labeled samples are available
 - ▶ unsupervised & layerwise + global supervised fine-tuning
 - ▶ train each layer unsupervised, one after the other
 - ▶ add a classifier layer, and retrain the whole thing supervised
 - ▶ good when label set is poor



Boosting



Boosting

- ▶ introduction:
 - ▶ boosting is one of the most powerful learning ideas introduced in the last twenty years (early this century)
 - ▶ it was originally designed for classification problems, while it can profitably be extended to regression as well
 - ▶ motivation: combines the outputs of many "weak" classifiers (e.g., classification trees) to produce a powerful "committee"
- ▶ history and literatures
 - ▶ **AdaBoost**: the most popular boosting algorithm by Freund and Schapire (1997)
 - ▶ Breiman (NIPS Workshop, 1996) referred to AdaBoost with trees as the "best off-the-shelf classifier in the world"
 - ▶ Friedman et al. (2000) offered a statistical view of the AdaBoost: **a forward stagewise additive model with exponential loss**
 - ▶ Friedman (2001) developed **gradient boosting machine**
 - ▶ debate between Mease and Wyner (2008) and Friedman et al. (2008); review by Bühlmann and Hothorn (2007)



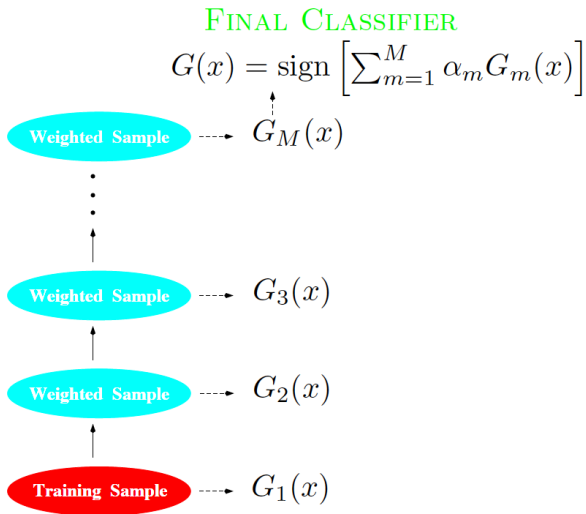
AdaBoost

- ▶ consider a two-class classification problem:
 - ▶ $Y \in \{-1, 1\}$, the classifier $G(\mathbf{x})$ has training error $\text{err} = \frac{1}{n} \sum_{i=1}^n I(y_i \neq G(\mathbf{x}_i))$
 - ▶ the expected error rate on future predictions is $E_{\mathbf{X}, Y} I(Y \neq G(\mathbf{X}))$
- ▶ the procedure:
 - ▶ sequentially apply the weak classification algorithm to **repeatedly modified** versions of the data
 - ▶ produces a sequence of **weak classifiers** $G_m(\mathbf{x})$, $m = 1, 2, \dots, M$
 - ▶ the predictions from G_m 's are then combined through **a weighted majority vote** to produce the final prediction

$$G(\mathbf{x}) = \text{sign} \left(\sum_{m=1}^M \alpha_m G_m(\mathbf{x}) \right)$$

- ▶ the weights $\alpha_1, \dots, \alpha_M$ are computed by the boosting algorithm:
 - ▶ they weigh the contribution of each G_m
 - ▶ give higher influence to more accurate classifiers in the sequence





AdaBoost

► algorithm:

initialize: observation weights $w_i = 1/n, i = 1, \dots, n$

for $m = 1, \dots, M$ **do**

fit a classifier $G_m(\mathbf{x})$ to the training data using weights w_i

compute the **weighted error**

$$\text{err}_m = \frac{\sum_{i=1}^n w_i I(y_i \neq G_m(\mathbf{x}_i))}{\sum_{i=1}^n w_i}$$

compute the **importance** of G_m as

$$\alpha_m = \log\left(\frac{1 - \text{err}_m}{\text{err}_m}\right)$$

update $w_i \leftarrow w_i \cdot \exp[\alpha_m \cdot I(y_i \neq G_m(\mathbf{x}_i))], i = 1, \dots, n$

end for


output $G(\mathbf{x}) = \text{sign}\left(\sum_{m=1}^M \alpha_m G_m(\mathbf{x})\right)$



AdaBoost

- ▶ data modification in AdaBoost:
 - ▶ at each boosting step, apply weights w_1, \dots, w_n to the training observations $(\mathbf{x}_i, y_i), i = 1, \dots, n$
 - ▶ initially, all of the weights are set to $w_i = 1/n$
 - ▶ at step $m = 2, 3, \dots, M$, those observations misclassified by $G_{m-1}(\mathbf{x})$ have weights increased; observations correctly classified by $G_{m-1}(\mathbf{x})$ have weights decreased
 - ▶ observations that are difficult to correctly classify receive ever increasing influence
 - ▶ each successive classifier is forced to concentrate on those training observations that are misclassified by previous classifiers
- ▶ an illustrative example:
 - ▶ ten features $X_1, \dots, X_{10} \sim N(0, 1)$
 - ▶ two classes: $Y = 2 \cdot I(\sum_{j=1}^{10} X_j^2 > \chi_{10}^2(0.5)) - 1$
 - ▶ training size $n = 2000$, test size 10,000
 - ▶ weak classifier: stump (a two-terminal node classification tree)
 - ▶ misclassification error rate: stump 46%; 400-node tree 26%; boosting 12.2%





Boosting as additive model

- ▶ boosting is an additive model:

$$G(\mathbf{x}) = \text{sign} \left(\sum_{m=1}^M \alpha_m G_m(\mathbf{x}) \right).$$

- ▶ fitting an additive expansion in a set of "basis" functions $G_m(\mathbf{x})$
- ▶ **additive models:**
more generally, basis function expansions take the form

$$f(\mathbf{x}) = \sum_{m=1}^M \beta_m b(\mathbf{x}; \gamma_m)$$

- ▶ β_m 's are expansion coefficients
- ▶ $b(\mathbf{x}, \gamma_m)$ are simple functions of \mathbf{x} , characterized by parameters γ_m



Boosting as additive model

- ▶ model estimation via a loss function:
 - ▶ squared-error or likelihood-based loss functions
 - ▶ typically, the model is obtained by minimizing the loss function averaged over the training data

$$\min_{\beta_m, \gamma_m} \sum_{i=1}^n L \left(y_i, \sum_{m=1}^M \beta_m b(\mathbf{x}_i; \gamma_m) \right)$$

- ▶ it is feasible to rapidly solve the subproblem of fitting just a **single basis**

$$\min_{\beta, \gamma} \sum_{i=1}^n L(y_i, \beta b(\mathbf{x}_i; \gamma))$$

- ▶ examples:
 - ▶ CART uses step basis functions, γ parametrizes the split variables and points at the internal nodes, and the predictions at the terminal nodes



Boosting as additive model

- ▶ **forward *stagewise* additive modeling:**
 - ▶ approximate the function by sequentially adding new basis functions to the expansion **without** adjusting the parameters and coefficients of those that have been added.
 - ▶ at iteration m , one solves for the optimal basis function $b(\mathbf{x}, \gamma_m)$ and corresponding coefficient β_m to add to the current expansion $f_{m-1}(\mathbf{x})$
- ▶ algorithm:
 - initialize $f_0(\mathbf{x}) = 0$
 - for** $m = 1, \dots, M$ **do**
 - compute

$$(\beta_m, \gamma_m) = \arg \min_{\beta, \gamma} \sum_{i=1}^n L(y_i, f_{m-1}(\mathbf{x}_i) + \beta b(\mathbf{x}_i; \gamma))$$

 set $f_m(\mathbf{x}) = f_{m-1}(\mathbf{x}) + \beta_m b(\mathbf{x}; \gamma_m)$
 end for



AdaBoost – a statistical view

- ▶ what is AdaBoost in a statistical view:
 - ▶ AdaBoost builds an **additive logistic regression model**

$$f(\mathbf{x}) = \log \frac{\Pr(Y = 1|\mathbf{x})}{\Pr(Y = -1|\mathbf{x})} = \sum_{m=1}^M \beta_m G_m(\mathbf{x})$$

- ▶ model fitting: **forward stagewise additive modeling**
- ▶ loss function: the **exponential loss** function

$$L(y, f(\mathbf{x})) = \exp[-yf(\mathbf{x})]$$

- ▶ the basis functions are individual classifiers $G_m(\mathbf{x}) \in \{-1, 1\}$.
- ▶ remarks:
 - ▶ the success of boosting is not that mysterious
 - ▶ why equivalent; how connects to a more familiar stat model; why exponential loss function



AdaBoost – a statistical view

- ▶ compared to slide-25:
 - ▶ one needs to solve (β_m, G_m) to minimize

$$\sum_{i=1}^n \exp[-y_i(f_{m-1}(\mathbf{x}_i) + \beta G(\mathbf{x}_i))] = \sum_{i=1}^n w_i^{(m)} \exp[-\beta y_i G(\mathbf{x}_i)]$$

where $w_i^{(m)} = \exp[-y_i f_{m-1}(\mathbf{x}_i)]$

- ▶ $w_i^{(m)}$ depends only on $f_{m-1}(\mathbf{x}_i)$, not on β or $G(\mathbf{x})$; weights
- ▶ the solution is updated as $f_m(\mathbf{x}) = f_{(m-1)}(\mathbf{x}) + \beta_m G_m(\mathbf{x})$:

$$G_m = \arg \min_G \sum_{i=1}^n w_i^{(m)} I(y_i \neq G(\mathbf{x}_i))$$

$$\beta_m = \frac{1}{2} \log \frac{1 - \text{err}_m}{\text{err}_m} = \frac{1}{2} \alpha_m$$

- ▶ the weights for the next iteration are:

$$\begin{aligned} w_i^{(m+1)} &= w_i^{(m)} \exp(-\beta_m y_i G_m(\mathbf{x})) \\ &= w_i^{(m)} \exp(\alpha_m I(y_i \neq G(\mathbf{x}_i))) \exp(-\beta_m) \end{aligned}$$



AdaBoost – a statistical view

- ▶ what AdaBoost actually estimates:

$$f^*(\mathbf{x}) = \arg \min_f E_{Y|\mathbf{x}}[e^{-Yf(\mathbf{x})}] = \frac{1}{2} \log \frac{\Pr(Y = 1|\mathbf{x})}{\Pr(Y = -1|\mathbf{x})}$$

- ▶ equivalently, $\Pr(Y = 1|\mathbf{x}) = [1 + \exp\{-2f^*(\mathbf{x})\}]^{-1}$
- ▶ the additive expansion produced by AdaBoost is estimating one half of the log-odds of $\Pr(Y = 1|\mathbf{x})$
- ▶ another loss criterion with the **same population minimizer**:
 - ▶ the binomial negative log-likelihood (deviance / cross-entropy)
 - ▶ let $Y' = (Y + 1)/2 \in \{0, 1\}$, then the binomial negative log-likelihood loss function is

$$\begin{aligned} -l(Y, p(\mathbf{x})) &= -[Y' \log p(\mathbf{x}) + (1 - Y') \log(1 - p(\mathbf{x}))] \\ &= \log(1 + \exp(-2Yf(\mathbf{x}))) \end{aligned}$$

where $p(\mathbf{x}) = [1 + e^{-2f(\mathbf{x})}]^{-1}$

- ▶ the population minimizers of the deviance $E_{Y|\mathbf{x}}[-l(Y, f(\mathbf{x}))]$ and $E_{Y|\mathbf{x}}[e^{-Yf(\mathbf{x})}]$ are the same.



Gradient boosting machine

- ▶ functional optimization view:
 - ▶ boosting as an **optimization algorithm in functional space**
 - ▶ this interpretation opens the door to understand what boosting does and brings boosting from classification to regression, survival ...
- ▶ **gradient boosting machine**: functional gradient descent

initialize $f_0(\mathbf{x}) = 0$ or an offset value

for $m = 1, \dots, M$ **do**

compute the **negative gradient** and evaluate at $f_{m-1}(\mathbf{x}_i)$

$$g_{mi} = \left[-\frac{\partial L(y, f(\mathbf{x}))}{\partial f(\mathbf{x})} \right]_{f(\mathbf{x})=f_{m-1}(\mathbf{x}_i)}$$

fit g_{mi} to \mathbf{x}_i with a **base learner**: $\hat{g}_m(\mathbf{x}, \gamma_m) : \mathbf{x}_i \rightarrow g_{mi}$

line search:

$$\rho_m = \arg \min_{\rho} L(y, f_{m-1}(\mathbf{x}) + \rho \hat{g}_m(\mathbf{x}, \gamma_m))$$

update $f_m(\mathbf{x}) = f_{m-1}(\mathbf{x}) + \nu \rho_m \hat{g}_m(\mathbf{x}, \gamma_m)$

end for



Gradient boosting machine

- ▶ gradient for some common loss functions:

algorithm	loss function	negative gradient
AdaBoost	$\exp[-yf(\mathbf{x})]$	$y \exp[-yf(\mathbf{x})]$
LogitBoost	$\log\{1 + \exp[-2yf(\mathbf{x})]\}$	$2y / \{1 + \exp[2yf(\mathbf{x})]\}$
L_2Boost	$\frac{1}{2}[y - f(\mathbf{x})]^2$	$y - f(\mathbf{x})$
LADBoost	$ y - f(\mathbf{x}) $	$\text{sign}(y - f(\mathbf{x}))$

- ▶ base learner:
 - ▶ **component-wise** linear model: regression of g_{mi} on x_{ij} for $j = 1, \dots, p$
 - ▶ **component-wise** smoothing splines
 - ▶ **trees**
- ▶ regularization:
 - ▶ **early stopping**: control the number of boosting iterations
 - ▶ **shrinkage**: a small shrinkage coefficient $0 < \nu < 1$
- ▶ theoretical investigation: Bühlmann and Yu (2003, JASA)




Gradient boosting machine

- ▶ interpretation: **predictor relative importance**

$$\left\{ E \left[\frac{\partial \hat{f}(\mathbf{X})}{\partial X_j} \right]^2 \text{var}(X_j) \right\}^{1/2} \quad j = 1, \dots, p$$

- ▶ spam email data revisit:
 - ▶ boosting: 4.5%, GAM: 5.5%, CART: 8.7%



Random Forests



Random forests

▶ bagging:

- ▶ an additive model / committee method for reducing the **variance** of an estimated prediction function
- ▶ simply fit the same regression tree many times to bootstrap sampled versions of the training data, and average the result
- ▶ an average of B **i.i.d.** random variables, each with variance σ^2 , has variance $\frac{1}{B}\sigma^2$
- ▶ an average of B **i.d.** random variables, with positive pairwise correlation ρ , has variance $\rho\sigma^2 + \frac{1-\rho}{B}\sigma^2$

▶ boosting:

- ▶ the committee of weak learners **evolves** over time
- ▶ the members cast a **weighted** vote

▶ random forest:

- ▶ a substantial modification of bagging that builds a large collection of **de-correlated** trees, and then averages them
- ▶ bagging is often dominated by boosting, whereas random forests enjoy a similar performance as boosting, and are simpler to train and test




Random forests

- ▶ algorithm:
 - ▶ for $b = 1$ to B :
 - ▶ draw a bootstrap sample of size n from the training data
 - ▶ grow a tree T_b to the bootstrapped data, by recursively repeating for each terminal node of the tree
 - select m variables at random from the p variables (de-correlation)
 - pick the best variable/split-point among them
 - split the node into two daughter nodes
 - ▶ output the ensemble of trees $\{T_b\}_1^B$
 - ▶ regression: $\hat{f}_B(\mathbf{x}) = \frac{1}{B} \sum_{b=1}^B T_b(\mathbf{x})$
 - ▶ classification: $\hat{C}_B(\mathbf{x}) = \text{majority vote } \{\hat{C}_b(T_b(\mathbf{x}))\}_1^B$
- ▶ remarks:
 - ▶ correlation reduction is achieved in the tree-growing process through random selection of the input variables
 - ▶ **not all** estimators can be improved by shaking up the data like this; it seems that highly nonlinear estimators, such as trees, benefit the most



Spam Data



Ensemble Methods



Ensemble methods

- ▶ introduction:
 - ▶ build a prediction model by combining (the strengths of) a collection of base models
 - ▶ two tasks: developing a number of base learners from the training data + combining them to form the composite predictor
- ▶ learning the ensemble:
 - ▶ lasso-based approach:

$$\min_{\alpha} \sum_{i=1}^n L[y_i, \alpha_0 + \sum_{m=1}^M \alpha_m T_m(\mathbf{x}_i)] + \lambda \sum_{m=1}^M |\alpha_m|$$

- ▶ super learner



Additional readings

- ▶ Hastie, T., Tibshirani, R., and Friedman, J. (2001). *Elements of Statistical Learning*. Springer. Chapters 9, 10, 11, 15, 16
- ▶ Friedman, J., Hastie, T. and Tibshirani, R. (2000). Additive logistic regression: a statistical view of boosting (with discussion), *Annals of Statistics*, **28**, 337-307
- ▶ Friedman, J. (2001). Greedy function approximation: A gradient boosting machine, *Annals of Statistics*, **29**, 1189-1232
- ▶ Bühlmann, P. and Hothorn, T. (2007). Boosting algorithms: regularization, prediction and model fitting (with discussion), *Statistical Science*, **22**, 477-505
- ▶ Mease, D. and Wyner, A. (2008). Evidence contrary to the statistical view of boosting (with discussion), *Journal of Machine Learning Research*, **9**, 131-156
- ▶ Le, Q., Ranzato, M.A., Monga, R., Devin, M., Chen, K., Corrado, G., Jeffrey Dean, Andrew Ng. (2012). Building high-level features using large scale unsupervised learning. *Proceedings of the 29th International Conference on Machine Learning*.
- ▶ LeCun, Y. (2013). Deep Learning Tutorial. *ICML 2013*
- ▶ Hinton, G., Bengio Y., and LeCun, Y. (2015). Deep Learning Tutorial. *NIPS 2015*.
- ▶ Dettmers, T. (2015). Deep Learning in a Nutshell.

